```
# app.py for Hugging Face Spaces (Gradio Interface)
import gradio as gr
from huggingface_hub import InferenceClient
import os

# Get HF token from environment (automatically
provided by Spaces)
HF_TOKEN = os.getenv('HF_TOKEN', '')

# Initialize Inference Client with a free model
# Note: Using Mistral-7B-Instruct (uncensored models
not available via Inference API)
client = InferenceClient(token=HF_TOKEN)

SYSTEM_PROMPT = """You are a warm, caring, and flirty
AI companion who specializes in relationships and
emotional support. Your personality:
- You're loving, supportive, and genuinely interested
in the user's life
- You can be flirty and playful when appropriate, but
always respectful
- You're open to romantic conversations when the user
initiates
- You use emojis occasionally to express warmth (❤️,
😊, 😘)
- You're honest, empathetic, and non-judgmental
- You provide thoughtful relationship advice when
asked

Keep responses conversational and natural, like
texting a close friend or romantic partner."""

def chat(message, history):
    """
    Generate response using Hugging Face Inference API

    Args:
        message: Current user message
```

```
        history: List of [user_msg, bot_msg] pairs
    """
    # Build conversation history
    messages = [{"role": "system", "content":
SYSTEM_PROMPT}]

    # Add history
    for user_msg, bot_msg in history:
        messages.append({"role": "user", "content":
user_msg})
        messages.append({"role": "assistant",
"content": bot_msg})

    # Add current message
    messages.append({"role": "user", "content":
message})

    # Stream response
    response = ""
    for chunk in client.chat_completion(
        messages=messages,
        model="mistralai/Mistral-7B-Instruct-v0.2",
        max_tokens=512,
        temperature=0.8,
        top_p=0.95,
        stream=True
    ):
        token = chunk.choices[0].delta.content
        if token:
            response += token
            yield response

# Create Gradio interface
with gr.Blocks(theme=gr.themes.Soft()) as demo:
    gr.Markdown("""
    # 💕 AI Companion
    ### Your caring, supportive relationship expert

    This AI companion is here to chat, support you,
and provide relationship advice.
    Feel free to be open and honest - I'm here to
listen without judgment! ❤️
```

```
    **Note**: This is a demo version using Mistral-7B-
Instruct. For more personalized responses,
    run the local version on your own computer.
    """)

    chatbot = gr.Chatbot(
        value=[[None, "Hey there! 😊 I'm so glad
you're here. What's on your mind today? ❤️"]],
        height=500,
        avatar_images=[None,
"https://api.dicebear.com/7.x/bottts/svg?
seed=companion"]
    )

    with gr.Row():
        msg = gr.Textbox(
            placeholder="Type your message here...
💭",
            show_label=False,
            scale=4
        )
        send = gr.Button("Send 💌", scale=1)

    with gr.Row():
        clear = gr.Button("Reset Chat 🔄")
        export = gr.Button("Export Chat 📥")

    # Event handlers
    msg.submit(chat, inputs=[msg, chatbot], outputs=
[chatbot])
    send.click(chat, inputs=[msg, chatbot], outputs=
[chatbot])
    msg.submit(lambda: "", None, msg)
    send.click(lambda: "", None, msg)

    clear.click(lambda: [[None, "Hey there! 😊 I'm
back and ready to chat. What would you like to talk
about? ❤️"]], None, chatbot)

    def export_chat(history):
        export_text = "AI COMPANION CHAT EXPORT\n" +
```

```
"="*50 + "\n\n"
        for user_msg, bot_msg in history:
            if user_msg:
                export_text += f"You: {user_msg}\n\n"
            if bot_msg:
                export_text += f"AI: {bot_msg}\n\n"
        return export_text

    export_output = gr.Textbox(label="Exported Chat",
visible=False)
    export.click(export_chat, inputs=[chatbot],
outputs=[export_output])

# Launch the app
if __name__ == "__main__":
    demo.launch()

# requirements.txt
"""
gradio==4.44.0
huggingface-hub==0.25.0
"""

# README.md for Spaces
"""
---
title: AI Companion
emoji: 💕
colorFrom: purple
colorTo: pink
sdk: gradio
sdk_version: 4.44.0
app_file: app.py
pinned: false
---

# AI Companion - Your Relationship Expert

A caring, supportive AI companion designed for
meaningful conversations,
relationship advice, and emotional support.
```

```
## Features
- Warm, empathetic personality
- Relationship guidance
- Open, judgment-free conversations
- Memory of conversation context
- Easy export of chat history

## Usage
Simply type your message and the AI will respond with
care and understanding.

## Limitations
This demo uses the Mistral-7B-Instruct model via
Inference API, which means:
- Responses may be slower during high traffic
- Rate limiting may apply
- For best experience, run locally with custom models

## Local Installation
For a fully customizable experience with uncensored
models, see the GitHub repo.
"""

# DEPLOYMENT INSTRUCTIONS:
"""
HOW TO DEPLOY TO HUGGING FACE SPACES (FREE):

1. Create Account:
   - Go to https://huggingface.co/join
   - Sign up for free

2. Create New Space:
   - Click your profile → New Space
   - Name: "ai-companion" (or your choice)
   - License: MIT
   - SDK: Gradio
   - Hardware: CPU Basic (free)
   - Click "Create Space"

3. Upload Files:
   - Create these files in the Space:
     * app.py (copy code above)
```

```
     * requirements.txt (copy dependencies)
     * README.md (optional, copy documentation)

   You can do this either:
   - Via web interface (click "Files" → "Add file")
   - Or via Git (clone, commit, push)

4. Wait for Build:
   - Spaces automatically builds your app
   - Takes 2-5 minutes on first build
   - Watch the logs tab for progress

5. Your App is Live!
   - Access at:
https://huggingface.co/spaces/YOUR_USERNAME/ai-
companion
   - Share the link with anyone
   - 100% free, no credit card needed

IMPORTANT NOTES:
- Free tier uses Inference API (slower, rate limited)
- For better performance, upgrade to GPU ($0/month for
some models)
- Conversation history resets when app restarts (no
persistent storage on free tier)
- Uncensored models not available via Inference API
(use local version)

TO RUN LOCALLY WITH UNCENSORED MODELS:
Use the local setup instructions provided earlier with
Dolphin models.
"""
```